Linnk Atlas AI v.1.0

Lead Generation & Staffing Sales Product Requirements Document (PRD)

Status: Draft for Product, Design, and Engineering Planning

Prepared by: Fazil A

Date: 31 May 2026

Introduction

This document defines requirements for Linnk Atlas AI v.1.0, an intelligence product designed to improve lead generation and outbound sales decision quality for staffing agencies.

Linnk Atlas AI is designed to answer one primary question: **How does Sarah (Sales Rep) know who deserves attention before outreach begins?**

This document intend to provide sufficient clarity for:

- Engineering teams to begin implementation planning
- Product Design teams to create workflows and interfaces
- Business stakeholders to evaluate investment decisions
- Product teams to assess scope, risks, and rollout readiness

Intended Audience

Primary Audience:

- Product
- Engineering
- Design
- Commercial Leadership
- Sales Operations

Secondary Audience:

- Agency Owners
- Sales Representatives
- Recruitment Teams

This PRD is intended to support:

- implementation planning
- prioritization decisions
- architecture discussions
- rollout planning
- launch readiness discussions

Sections are intentionally structured to move from: **problem definition → product decisions → system design → operationalization.**

Executive Summary

Problem Statement

Staffing sales representatives do not primarily struggle because they cannot make calls. They struggle because they do not know which calls deserve attention today.

Sarah begins her morning with:

- hundreds of possible companies
- fragmented information spread across CRM, inboxes, spreadsheets, LinkedIn, ATS systems, and personal notes
- limited visibility into hiring demand signals
- significant manual research effort
- uncertainty about which opportunities are worth prioritizing

As a result:

- large portions of outreach target weak opportunities
- strong opportunities are discovered too late
- sales productivity depends heavily on instinct rather than evidence
- representatives spend significant time searching before selling

Existing CRMs primarily record activity after work occurs. They do not substantially improve decision quality before work begins.

Product Thesis

Linnk Atlas AI is an intelligence layer built for staffing sales teams. Atlas operates alongside existing CRM, ATS, inbox, and workflow systems rather than replacing them.

Link Atlas AI helps representatives:

- identify who deserves attention today
- understand why the opportunity matters now
- reduce time spent identifying and validating opportunities
- prepare for conversations faster

Atlas is designed to influence sales decisions before outreach occurs rather than documenting activity afterward.

Core hypothesis: **If high-signal opportunities are surfaced earlier and decision friction is reduced, then representatives spend less time searching and more time selling.**

Business Goal

Increase revenue productivity of staffing sales teams without proportional increases in sales headcount. The product should enable agencies to:

- increase sales output per representative
- reduce time spent researching opportunities
- increase qualified meetings generated per representative
- increase meeting generation beyond the current baseline of approximately three meetings per week

Product Goal

Enable Sarah to open Linnk Atlas AI each morning and determine within minutes rather than hours:

1. **Who should I contact today?**
2. **Why does this opportunity matter now?**
3. **How should I prepare before speaking with them?**

Success requires changing Sarah's morning workflow from: **Research** → **Decide** → **Sell** to **Review** → **Decide** → **Sell**

Problem Definition

Current Workflow (Sarah Today)

Sarah's morning workflow today:

1. Open CRM
2. Open spreadsheets
3. Search LinkedIn / Job Boards
4. Review emails and notes
5. Guess which accounts deserve attention
6. Begin outreach
7. Update systems later

This workflow creates substantial decision friction before selling activity begins.

Primary Problems

Problem 1 — Weak Opportunity Prioritization

Sarah starts each day with hundreds of possible accounts but limited confidence regarding:

- who deserves attention today
- why timing matters now
- which opportunities are strengthening
- which opportunities are weakening

As a result:

- weak opportunities consume time

- strong opportunities are missed
- productivity depends heavily on instinct

Result: **High activity volume with inconsistent decision quality**

Problem 2 — Research Overhead

Before outreach begins, significant time is spent gathering information from:

- company websites
- job boards
- emails
- CRM records
- previous conversations

Time spent researching directly reduces time available for selling.

Result: **Representatives spend substantial time searching before acting**

Problem 3 — Fragmented Signals

Signals required for prioritization exist across multiple disconnected systems.

External Sources:

- hiring activity
- leadership changes
- company announcements
- expansion signals
- tenders

Internal Sources:

- CRM
- ATS
- inboxes
- spreadsheets
- previous interactions

No system consolidates these signals into actionable decisions.

Result: **Important opportunities remain operationally invisible**

Problem 4 — Low Trust In Existing Workflows

Existing CRM systems primarily function as **systems of record** rather than **systems that influence decisions**

Observed behaviors:

- spreadsheet dependence
- CRM avoidance
- shadow workflows

Result: **Representatives work outside existing systems**

Secondary Problems

Problem 1 — Administrative Burden

After outreach activity, representatives still spend time:

- updating CRM
- summarizing calls
- recording notes
- creating follow-ups

This work contributes to lower selling time but is considered secondary to decision quality problems.

Jobs To Be Done

1. Functional Job: Help Sarah determine where my selling time creates the highest return today.
2. Emotional Job: Reduce uncertainty that Sarah is spending effort in the wrong places.
3. Social Job: Help Sarah perform consistently without relying entirely on instinct and personal memory.

Existing Alternatives

Today representatives compensate using:

- CRM dashboards
- spreadsheets
- manual research
- experience
- manager guidance

These approaches are:

- inconsistent
- difficult to scale
- heavily dependent on individuals

Success Criteria

The product successfully solves the problem when:

- time spent identifying opportunities reduces from hours to minutes
- representatives consistently act on surfaced opportunities
- time spent selling increases relative to research time
- meeting generation improves beyond current baseline performance
- sales productivity improves without proportional headcount growth

Users & Personas

This product is designed for staffing organizations. However, the product is intentionally optimized around one primary user first: **Sarah**

- The product does not attempt to solve staffing operations broadly.
- It is designed primarily to improve decision quality for outbound staffing sales workflows.
- This scope decision constrains product priorities throughout this document.

Primary Persona: Sarah

- **Role:** Business Development Representative
- **Experience:** 10+ years staffing sales experience
- **Agency Size:** ~60 employees
- **Domain:** Engineering staffing for manufacturing and energy clients
- **Primary Goal:** Generate more qualified meetings without increasing working hours.
- **Current Reality:** Sarah manages large numbers of potential accounts with fragmented information and limited visibility into which opportunities deserve attention. Her success depends heavily on making correct prioritization decisions.

How Sarah Measures Success

1. Sarah is primarily evaluated on:
    a. meetings booked
    b. opportunities generated
    c. pipeline growth
    d. placements influenced
    e. client relationships
2. Sarah is NOT evaluated on:
    a. amount of research performed
    b. CRM completeness

    c. number of systems used

Implication: **The product must improve selling outcomes rather than create additional workflow burden.**

Sarah's Constraints

1. Time Constraint: Research competes directly with selling activity. Every additional minute spent searching reduces time available for calls and relationship building.
2. Information Constraint: Important signals exist. They remain fragmented across:
    a. Internal Systems:
        i. CRM
        ii. ATS
        iii. Email
    b. External Sources:
        i. Hiring activity
        ii. Leadership changes
        iii. Company announcements
        iv. Market signals
3. Trust Constraint: Sarah possesses significant domain experience and existing workflows.
    a. Recommendations must therefore be:
        i. evidence backed
        ii. explainable
        iii. Actionable
4. Cognitive Constraint: Sarah manages large numbers of accounts and possible actions simultaneously. The product must reduce decision load rather than increase it.

Secondary Personas

Secondary Persona A — Sales Team Lead

Primary Goal: Increase team productivity and improve outbound effectiveness.

Questions:

- Are recommendations reliable?
- Is team productivity improving?
- Are representatives adopting workflows?

Primary Requirement: Visibility into usage and recommendation quality.

Secondary Persona B — Agency Owner

Primary Goal: Increase revenue productivity without proportional hiring growth.

Questions:

- Does output per representative improve?
- Is measurable ROI visible?
- Does this reduce dependency on headcount growth?

Primary Requirement: Clear productivity measurement.

Anti Personas

1. Recruiters Focused Primarily On Candidate Delivery

    Reason: Primary problem being solved is sales prioritization. Not candidate sourcing.

2. Organizations Seeking Immediate CRM Replacement

    Reason: Replacement creates

    - higher implementation friction
    - longer deployment cycles
    - greater change management burden

Persona Prioritization

Priority Order:

1. Sarah
2. Sales Team Lead
3. Agency Owner

If conflicts emerge **Sarah wins.**

Reason: Without Sarah adoption, downstream value creation does not occur.

Product Principles

These principles define constraints under which Linnk Atlas AI will be built. They exist to guide:

- product decisions
- engineering tradeoffs
- AI usage decisions
- prioritization decisions
- launch decisions

If future decisions conflict with these principles, these principles take precedence unless explicitly revised.

Principle 1 — Evidence Before Automation

Recommendations should be understandable before they are trusted. Requirements:

- recommendations must expose evidence
- signal sources should remain visible
- timing context should remain visible
- recommendations should explain why action is suggested

Example:

Poor: **Call Acme Manufacturing**

Better: **Call Acme Manufacturing**

Why:

- 4 engineering roles posted within 10 days
- manufacturing expansion announced
- previous relationship exists

Principle 2 — Assist Humans Before Replacing Humans

The product should reduce cognitive work before attempting to remove human decisions.

Autonomy Ladder:

1. Suggest
2. Prepare
3. Assist
4. Act with approval
5. Act autonomously

Target Autonomy: **Suggest → Prepare → Assist**

Allowed:

- recommendations
- summaries
- outreach preparation
- workflow assistance

Not Allowed:

- autonomous outreach
- autonomous account decisions
- customer communication without approval

Principle 3 — Integrate Existing Workflows Before Creating New Ones

Linnk Atlas AI should operate alongside existing systems rather than replace them.

Requirements:

- CRM integration
- ATS integration
- inbox integration
- export capabilities

Avoid:

- mandatory migrations
- workflow replacement

Principle 4 — AI Requires Explicit Justification

AI is not the default decision. Every AI capability must answer:

1. Why AI here?
2. What does the non-AI version look like?
3. What measurable advantage does AI create?

Examples:

1. Signal Collection:
    a. Decision: **No AI**
    b. Reason: Structured ingestion is simpler and more reliable.
2. Signal Interpretation:
    a. Decision: **AI**
    b. Reason: Large volumes of unstructured information require interpretation.
3. Lead Prioritization:
    a. Decision: **Hybrid**
    b. Reason: Pure rules are insufficient while pure AI is difficult to trust.

Principle 5 — Design For Failure

Recommendations will sometimes be incorrect. The product must therefore support:

- monitoring
- fallback behavior
- feedback loops
- rollback mechanisms
- evaluation systems

Principle Validation Test

Features should pass the following checks:

- Does it improve decision quality?
- Can users understand why it exists?
- Can users recover when wrong?
- Does AI create measurable advantage?
- Does it reduce work rather than create work?
- Does it preserve trust?

If any answer is no: **Feature requires re-evaluation**

Scope Definition

Linnk Atlas AI is intentionally scoped as a sales intelligence layer that improves outbound sales decision quality before outreach begins.

Atlas is not intended to become a complete staffing operating system.

1. **Primary problem solved:** How does Sarah know who deserves attention before outreach begins?
2. **Secondary problem solved:**
    a. How can Sarah prepare and execute faster once opportunities are identified?
    b. How can important follow-ups avoid becoming invisible?

Atlas does not primarily solve:

- candidate delivery
- payroll
- recruitment operations
- CRM replacement

The scope intentionally prioritizes improving decision quality before optimizing downstream workflows.

Linnk Atlas AI MVP

Capability 1 — Opportunity Intelligence

Purpose: Identify which companies deserve attention today and explain why.

Capabilities

- ingest external signals
- consolidate fragmented information
- identify emerging opportunities
- rank opportunities
- explain recommendation rationale

Outputs

- ranked opportunities
- evidence visibility
- recommendation rationale

AI Decision: AI Used

Why AI Here?

Opportunity identification requires interpreting large volumes of partially structured and unstructured signals across:

- hiring pages
- job postings
- leadership changes
- company announcements
- company websites
- news signals

Deterministic systems can collect signals.

They struggle to:

- interpret relationships across multiple weak signals
- summarize noisy information into actionable reasoning
- generalize across new signal patterns

AI creates value through interpretation rather than collection.

Non-AI Alternative

Build:

- deterministic ingestion pipelines
- weighted rule engines
- keyword matching systems
- manually configured scoring models

What Still Works Without AI?

- signal collection
- company ranking
- basic prioritization
- workflow creation

What Is Lost Without AI?

- contextual interpretation
- synthesized reasoning
- adaptability to new signal patterns
- ability to combine weak signals into stronger opportunity hypotheses

Capability 2 — Follow-Up Intelligence

Purpose: Ensure important opportunities do not disappear after initial interaction.

Capabilities:

- aging opportunity detection
- follow-up prioritization
- re-engagement recommendations
- urgent follow-up queues

Outputs:

- prioritized follow-up queue
- recommended actions
- urgency visibility

AI Decision: Hybrid

Why AI Here?

Rules can detect inactivity. AI creates value when determining:

- whether opportunities remain valuable
- whether new signals justify re-engagement
- recommended timing and approach

Non-AI Alternative

Build:

- reminder systems
- activity timers
- workflow rules

What Still Works Without AI?

- reminders
- activity queues
- follow-up scheduling

What Is Lost Without AI?

- contextual follow-up prioritization
- adaptive timing

- signal-aware re-engagement

Capability 3 — Decision Support Workspace

Purpose: Provide operational workflow for daily prioritization.

Capabilities

- prioritized opportunity feed
- follow-up queues
- urgency indicators
- daily workflow management

Outputs

- daily working environment
- visible prioritization
- actionable workflows

AI Decision: No AI

Why Not AI Here?

- Interfaces and workflow management are deterministic problems.
- AI operates behind recommendations rather than inside the workspace itself.

Capability 4 — Outreach Preparation Assistance

Purpose: Reduce preparation effort before calls or outreach.

Capabilities

- summarize opportunity context
- surface relevant historical context
- suggest conversation angles

Outputs

- preparation support
- reduced research effort

AI Decision: AI Used

Why AI Here?

Preparation requires compressing fragmented information into usable context. AI creates value through:

- summarization
- information compression
- contextual synthesis

Deterministic systems struggle when information sources become large, fragmented, and inconsistent.

Non-AI Alternative

Build:

- templates
- static account notes
- manually maintained account summaries
- manual research workflows

What Still Works Without AI?

- account history visibility
- templates
- preparation workflows

What Is Lost Without AI?

- personalized preparation
- faster context creation
- dynamic synthesis across multiple sources
- reduced preparation time

Capability 5 — Administrative Assistance

Purpose: Reduce post-activity workload.

Capabilities

- summarize calls
- suggest follow-ups
- recommend CRM updates

Outputs

- reduced administrative burden
- suggested actions

AI Decision: AI Used

Why AI Here?

Activity data is conversational and unstructured. AI enables:

- summarization
- extraction of actions
- structured suggestions from messy inputs

Non-AI Alternative

Build:

- structured activity forms
- templates
- workflow automation rules
- manual updates

What Still Works Without AI?

- CRM workflows
- manual activity capture
- template driven updates

What Is Lost Without AI?

- reduced manual effort
- automatic summarization
- faster post-call workflows

Capability 6 — Feedback & Learning Loops

Purpose: Enable evaluation and recommendation improvement.

Capabilities

- feedback capture
- corrections
- dismissals
- monitoring signals

AI Decision: No AI

Why Not AI Here?

Feedback collection is deterministic and should remain explicit.

Explicit MVP Scope

Version 1 Includes:

- Opportunity Intelligence
- Follow-Up Intelligence
- Linnk Workspace
- Preparation Assistance
- CRM Assistance
- Feedback Loops

Version 1 Excludes:

- candidate workflows
- autonomous communication
- CRM replacement
- advanced forecasting

Explicitly Out Of Scope

1. Full CRM Replacement

   Reason:

   - high switching cost
   - implementation complexity
   - slower adoption

   Decision: Integrate existing systems.

2. Candidate Matching & Candidate Search

   Reason: Primary problem is identifying opportunities rather than filling opportunities.

3. Autonomous Sales Agent

   Excluded:

   - automatic emails
   - automatic messaging
   - autonomous sequencing

   Reason: Relationship and trust risk exceeds productivity benefit.

   Decision: Human approval required.

AI Explicitly Rejected

1. External Signal Collection

   Decision: **No AI**

   Reason: Signal collection is primarily structured ingestion.

   Deterministic pipelines are:

   - cheaper
   - simpler
   - more reliable

2. Final Opportunity Decision

   Decision: **Human decides**

   Reason: Opportunity selection directly impacts relationships, trust, and sales performance.

3. Automatic CRM Modification

   Decision: **No automatic updates**

   Reason: Incorrect records create operational risk. Users approve changes before submission.

Scope Prioritization Logic

Priority Order:

1. Opportunity Intelligence
2. Follow-Up Intelligence
3. Decision Support Workspace
4. Preparation Assistance
5. Administrative Reduction

Reason: Improving prioritization creates larger value than optimizing downstream workflows.

Definition Of Scope Success

Scope succeeds when:

- Sarah opens Linnk Atlas AI before CRM
- recommendations influence behavior
- follow-ups happen more consistently
- research time reduces
- trust remains sufficiently high for continued usage
- existing systems remain systems of record

Product Overview

Linnk Atlas is designed as a sales intelligence layer operating above existing staffing workflows. Atlas changes what happens before sales activity begins and continues supporting work until follow-ups are completed.

The product is designed around one behavioral objective: **Sarah should know where to spend her selling time within minutes of opening the product.**

Operational Product Flow

1. Signals Collected
2. Signals Interpreted
3. Opportunities Generated
4. Linnk Workspace Created
5. Sarah Reviews Priorities
6. Sarah Prepares Outreach
7. Sales Activity Occurs
8. Follow-Up Priorities Generated
9. Feedback Captured
10. Recommendations Improve

Product Architecture Overview

Linnk Atlas AI consists of seven operational layers.

Layer 1 — Data Sources Layer

- Responsibility: Collect internal and external signals.

- Output: Raw signals.

Layer 2 — Signal Intelligence Layer

- Responsibility: Convert fragmented information into structured signals.
- Output: Structured evidence.

Layer 3 — Opportunity Intelligence Layer

- Responsibility: Generate ranked opportunities.
- Output: Prioritized opportunities.

Layer 4 — Linnk Workspace

- Responsibility: Create operational workflow.
- Output: Daily priorities.

Layer 5 — Outreach Preparation Layer

- Responsibility: Reduce research burden.
- Output: Preparation support.

Layer 6 — Workflow Assistance Layer

- Responsibility: Support post-outreach work and follow-up prioritization.
- Outputs:
  - suggested updates
  - follow-up recommendations
  - administrative assistance

Layer 7 — Feedback & Evaluation Layer

- Responsibility: Measure recommendation quality and support iteration.
- Output: Evaluation signals.

Product Autonomy Summary

Linnk Atlas AI intentionally targets:

1. Interpret
2. Recommend
3. Prepare
4. Assist

Linnk Atlas AI intentionally avoids:

1. Act autonomously
2. Operate independently
3. Relationship ownership remains human.

Detailed User Journeys

This section defines primary workflows users experience while using Linnk Atlas AI.

These workflows are intended to support:

- Product Design
- Engineering
- QA
- Product Development

Each workflow defines:

- trigger
- workflow
- system behavior

- success state
- failure handling

Journey 1 — Linnk Prioritization Workflow

Trigger: Sarah opens Atlas at the start of her day.

Workflow

Step 1: Sarah opens Atlas.

Step 2: Atlas displays:

- Top Opportunities Today
- Urgent Follow-Ups
- Emerging Accounts
- Recently Active Accounts

Step 3: Sarah reviews recommendations.

Each recommendation displays:

- company
- priority
- why surfaced
- evidence
- signal recency
- recommended next action

Step 4: Sarah chooses to

- open recommendation
- prepare outreach
- dismiss recommendation
- save recommendation
- mark recommendation incorrect

System Behavior: Linnk Atlas AI

- loads latest signals
- generates rankings
- applies business rules
- generates explanations

Success State: Sarah identifies

- who deserves attention
- why they deserve attention

within approximately five minutes.

Failure Handling: If recommendations appear irrelevant

- broader opportunity pool displayed
- dismissals captured
- feedback collected

Journey 2 — Opportunity Review Workflow

Trigger: Sarah selects an opportunity.

Workflow

Atlas displays

- company overview
- supporting signals
- previous interactions

- recommended actions

Sarah chooses:

- prepare outreach
- open CRM
- save opportunity
- dismiss opportunity

Success State: Sarah understands

- why opportunity surfaced
- whether action is justified

Failure Handling: If signals appear weak

- dismissal allowed
- corrections allowed
- feedback captured

Journey 3 — Outreach Preparation Workflow

Trigger: Sarah selects **Prepare Outreach**

Workflow

Linnk Atlas generates:

- summarized context
- previous interactions
- suggested talking points
- conversation angles

Sarah:

- edits suggestions
- copies suggestions
- ignores suggestions

Success State: Sarah prepares faster than existing workflow.

Failure Handling: If suggestions appear poor

- regeneration available
- editing supported
- feedback collected

Journey 4 — Follow-Up Intelligence Workflow

Trigger: Existing opportunities remain inactive beyond expected periods OR new signals appear.

Workflow

Atlas surfaces:

- aging opportunities
- urgent follow-ups
- re-engagement suggestions

Sarah chooses:

- contact now
- schedule later
- dismiss recommendation

System Behavior: Atlas combines

- activity history
- signal recency

- relationship history

Success State: Important opportunities do not disappear after initial outreach.

Failure Handling: If follow-up recommendation appears poor:

- dismissal supported
- feedback collected

Journey 5 — Activity Capture Workflow

Trigger:

- Call completed.
- Email sent.
- Meeting completed.

Workflow

Linnk Atlas generates:

- activity summary
- suggested follow-ups
- suggested CRM updates

Sarah:

- approves
- edits
- rejects

Product Rule: Atlas never automatically

- updates CRM
- sends communication
- modifies customer records

  User approval required.

Success State: Activity logging becomes faster.

Failure Handling: Incorrect summaries

- editable
- rejectable
- corrections captured

Journey 6 — Recommendation Feedback Workflow

Trigger: User interacts with recommendations.

Available Actions: Users can:

- act
- dismiss
- save
- mark incorrect
- provide feedback

System Behavior: Atlas stores:

- feedback events
- recommendation metadata
- corrections

Success State: Sufficient feedback exists for evaluation.

Journey 7 — Low Confidence Workflow

Trigger: Recommendation confidence below threshold.

System Behavior:

Instead of: **Call Company X**

Atlas displays: **Potential Opportunity**

Atlas additionally displays:

- supporting evidence
- confidence indicators
- alternatives

Product Rule: Low confidence recommendations

- receive weaker prioritization
- never trigger automation

Success State: Users understand uncertainty rather than receiving false certainty.

User Journey Success Definition

Workflows succeed when:

- Sarah opens Atlas before CRM
- recommendations influence behavior
- research effort decreases
- follow-ups happen consistently
- administrative effort decreases
- trust improves over time

Functional Requirements

This section defines functional requirements required for **Link Atlas v.1.0** MVP launch.

Requirements are grouped by capability. Requirement format:

- Requirement ID
- User Story
- Requirement Description
- Inputs
- Outputs
- Priority
- Dependencies
- Acceptance Criteria

Shared Objects

Signal Object: Required Fields

- Signal ID
- Company ID
- Source
- Signal Type
- Timestamp
- Evidence
- Confidence Band

Opportunity Object: Required Fields

- Opportunity ID
- Company ID
- Priority
- Explanation
- Supporting Signals

- Recommended Action
- Confidence Band

Recommendation Object: Required Fields

- Recommendation ID
- Recommendation Type
- Timestamp
- Supporting Evidence
- User Actions

Feedback Object: Required Fields

- Feedback ID
- Recommendation ID
- Feedback Type
  - User Action
  - Correction Value
  - Optional Comment
- Timestamp
- User ID
- Model Version
- Context Metadata
- Supported Feedback Types:
  - Accept
  - Dismiss
  - Save
  - Incorrect
  - Override
  - Comment

Follow-Up Object: Required Fields

- Follow-Up ID
- Company ID
- Associated Recommendation ID
- Follow-Up Type
- Suggested Action
- Due Date
- Priority
- Reason
- Status
- Confidence
- Created Timestamp

Opportunity Intelligence Requirements

FR-OI-01 — Signal Ingestion

User Story: Sarah should not manually search multiple systems every morning to identify opportunities.

Requirement Description: System continuously ingests supported internal and external signals.

Inputs

Internal:

- CRM activity
- ATS records
- previous interactions

External:

- job boards
- company websites
- career pages
- news

- leadership changes
- tenders

Outputs: Normalized raw signals.

Priority: P0

Dependencies: Integrations Layer, Ingestion Infrastructure

Acceptance Criteria

- 95% ingestion success rate
- ingestion failures generate alerts
- failed sources visible internally

FR-OI-02 — Signal Normalization

User Story: Sarah should receive standardized recommendations regardless of source complexity.

Requirement Description: Convert raw signals into standardized Signal Objects.

Inputs: Raw ingested signals.

Outputs: Signal Objects.

Priority: P0

Dependencies: Signal Ingestion

Acceptance Criteria

- duplicate signals consolidated
- signals normalized before ranking

FR-OI-03 — Opportunity Generation

User Story: Sarah should receive prioritized opportunities rather than raw information.

Requirement Description: Generate Opportunity Objects from normalized signals.

Inputs

- Signal Objects
- relationship history
- CRM history

Outputs: Opportunity Objects.

Priority: P0

Dependencies: Signal Normalization

Acceptance Criteria

- opportunities generated for >90% supported accounts
- every opportunity contains explanation

FR-OI-04 — Opportunity Ranking

User Story: Sarah should understand which opportunities deserve immediate attention.

Requirement Description: Rank opportunities using signal strength, relationship context and business constraints.

Inputs

- opportunity objects
- business rules
- relationship history

Outputs: Ranked opportunities.

Priority: P0

Dependencies: Opportunity Generation

Acceptance Criteria

- ranking refresh <4 hours
- ranking generation <5 seconds P95
- low confidence recommendations excluded from Top 5
- maximum displayed opportunities = 20

Follow-Up Intelligence Requirements

FR-FI-01 — Follow-Up Queue Generation

User Story: Sarah should not forget important opportunities after first interaction.

Requirement Description: Generate prioritized follow-up queues.

Inputs

- previous activity
- relationship history
- signal recency

Outputs

- urgent follow-ups
- aging opportunities
- re-engagement suggestions

Priority: P0

Dependencies: Opportunity Intelligence

Acceptance Criteria

- queue refreshed daily
- inactive opportunities surfaced automatically

FR-FI-02 — Follow-Up Prioritization

User Story: Sarah should know which follow-ups deserve attention first.

Requirement Description: Prioritize follow-ups using activity history and opportunity context.

Inputs

- inactivity duration
- new signals
- engagement history

Outputs: Prioritized follow-up queue.

Priority: P0

Dependencies: Follow-Up Queue Generation

Acceptance Criteria

- duplicate follow-ups prevented
- refresh occurs daily

Linnk Workspace Requirements

FR-MW-01 — Workspace Generation

User Story: Sarah should understand where to spend time within minutes.

Requirement Description: Generate daily workspace after recommendations are produced.

Outputs

- Top Opportunities
- Urgent Follow-Ups
- Emerging Accounts
- Recently Active Accounts

Priority: P0

Dependencies: Ranking Engine, Follow-Up Intelligence

Acceptance Criteria

- workspace loads <5 seconds P95

FR-MW-02 — Opportunity Cards

Requirement Description: Each opportunity card displays required recommendation information.

Required Fields

- company
- explanation
- confidence
- signal recency
- recommended action

Available Actions

- open
- save
- dismiss
- prepare outreach
- open CRM

Priority: P0

Acceptance Criteria

- 100% displayed recommendations include explanations

Outreach Preparation Requirements

FR-OP-01 — Preparation Package Generation

User Story: Sarah should prepare for outreach faster.

Requirement Description: Generate preparation packages before outreach begins.

Inputs

- opportunity object
- relationship history
- previous interactions

Outputs

- summary
- signals

- relationship context
- conversation angles

Priority: P0

Dependencies: Opportunity Intelligence

Acceptance Criteria

- package generation <5 seconds P95
- missing inputs handled gracefully

FR-OP-02 — Outreach Support

Requirement Description: Generate optional outreach support.

Outputs

- email starters
- message starters
- call openers

Priority: P1

Constraints

- editable only
- no auto-send
- no automatic storage

Acceptance Criteria

- generated outputs always editable

Workflow Assistance Requirements

FR-WA-01 — Activity Summaries

User Story: Sarah should spend less time updating systems after outreach.

Requirement Description: Generate activity summaries.

Inputs

- calls
- notes
- meetings

Outputs

- summaries
- suggested follow-ups
- suggested CRM updates

Priority: P1

Acceptance Criteria

- availability >95%

FR-WA-02 — Approval Workflow

Requirement Description: Users approve all system modifications.

Priority: P0

Acceptance Criteria

- CRM modifications require approval

- communication never automatically sent
- records never modified automatically

Feedback Requirements

FR-FB-01 — Feedback Capture

Priority: P0

Requirement Description: Capture recommendation feedback.

Feedback Types

- accept
- dismiss
- save
- incorrect
- comments

Acceptance Criteria

- feedback submitted within <=1 interaction

FR-FB-02 — Recommendation Traceability

Priority: P0

Requirement Description: Store recommendation lineage.

Acceptance Criteria

- 100% recommendations traceable

Non Functional Requirements

Performance

- workspace load <5 seconds P95
- recommendation generation <5 seconds P95
- preparation generation <5 seconds P95

Availability

- 99.5%

Security

- RBAC
- encryption
- audit logging

Explainability: Every recommendation exposes

- evidence
- rationale
- signal recency

Requirements Success Definition

- engineering can estimate work
- QA can derive tests
- design can define interfaces
- launch gates become measurable

System Architecture

This section defines major system components, integration boundaries, operational dependencies, and data flow required to support Linnk Atlas.

This section answers:

- What components exist?
- How components interact?
- Where AI operates?
- How failures are handled?
- What dependencies exist?

The objective is implementation clarity rather than infrastructure specification.

High-Level Architecture

1. Internal + External Sources (data sources)
2. Integration Layer
3. Signal Processing Layer
4. Opportunity Intelligence Layer
5. Recommendation Services
6. Linnk Atlas Workspace
7. User Actions
8. Feedback & Evaluation Layer

Component Architecture

Component 1 — Integration Layer

Purpose: Collect information from connected systems.

Responsibilities

- retrieve data
- synchronize sources
- validate connectivity
- maintain source health

Inputs

Internal Sources:

- CRM
- ATS
- email metadata
- previous activity

External Sources:

- job boards
- company websites
- hiring pages
- leadership changes
- news
- tenders

Outputs: Raw Signal Objects

Dependencies

- source integrations
- authentication
- permissions

Component 2 — Signal Processing Layer

Purpose: Convert raw signals into standardized inputs.

Responsibilities

- normalization
- deduplication
- classification
- evidence creation

Inputs: Raw signals.

Outputs: Signal Objects.

Dependencies: Integration Layer.

Component 3 — Opportunity Intelligence Layer

Purpose: Generate prioritized opportunities.

Responsibilities

- opportunity generation
- prioritization
- ranking
- follow-up prioritization

Inputs

- signal objects
- relationship history
- previous activity

Outputs: Opportunity Objects.

Dependencies: Signal Processing Layer.

Component 4 — Recommendation Services

Purpose: Generate actionable recommendations.

Responsibilities

- explanations
- recommendation generation
- confidence scoring
- preparation generation

Inputs: Opportunity Objects.

Outputs: Recommendation Objects.

Dependencies: Opportunity Intelligence Layer.

Component 5 — Linnk Atlas Workspace

Purpose: Provide operational workflow environment.

Responsibilities

- display recommendations
- capture actions
- trigger workflows

Outputs: User interactions.

Dependencies: Recommendation Services.

Component 6 — Feedback & Evaluation Layer

Purpose: Measure quality and enable improvement.

Responsibilities

- feedback collection
- recommendation evaluation
- monitoring
- quality measurement

Outputs: Evaluation signals.

Dependencies: User interactions.

Data Flow

1. Step 1: Signals collected through integrations.
2. Step 2: Signals normalized and structured.
3. Step 3: Opportunities generated.
4. Step 4: Recommendations produced.
5. Step 5: Workspace generated.
6. Step 6: Users interact with recommendations.
7. Step 7: Feedback captured.
8. Step 8: Evaluation pipelines updated.

AI Placement Architecture

AI operates only inside components requiring interpretation.

AI Components

- Signal Processing Layer
- Opportunity Intelligence Layer
- Recommendation Services
- Workflow Assistance

Non-AI Components

- Integrations
- Workspace orchestration
- Monitoring
- Feedback collection

Reason: Deterministic systems preferred when measurable advantage from AI is weak.

Integration Model

Atlas operates alongside existing systems. Atlas integrates with

- CRM systems
- ATS systems
- inbox systems

Atlas does not become system of record. System of record ownership remains external.

Failure Handling Architecture

- Failure: Signal Source Unavailable
    - Response: Use cached signals and ingestion alerts created
- Failure: Recommendation Engine Failure
    - Response: Fallback ranking

- Failure: AI Failure
  - Response: Fallback deterministic workflows and explanations exposed
- Failure: Integration Failure
  - Response: Continue operating with partial functionality and affected sources surfaced

Operational Requirements

Performance

- signal freshness <4 hours
- workspace load <5 seconds P95
- recommendation generation <5 seconds P95

Availability

- target availability 99.5%

Security

- RBAC
- encrypted storage
- audit logging
- secure integrations

AI Specification

This section defines how AI systems operate inside Linnk Atlas including:

- responsibilities
- boundaries
- autonomy levels
- inputs / outputs
- evaluation requirements
- failure handling

This section answers:

- Where does AI belong?
- Where does AI explicitly not belong?

Atlas follows one rule: **AI exists only where deterministic approaches create meaningful limitations.**

AI Operating Principles

Principle 1 — Assist Before Acting

Target Autonomy:

1. Suggest
2. Prepare
3. Assist

Atlas intentionally avoids:

- autonomous execution
- independent operation

Principle 2 — Evaluation Before Models

- Models may change.
- Evaluation standards should remain stable.
- Model selection follows evaluation requirements rather than the reverse.

Principle 3 — Structured Outputs Over Free Text

Model outputs should follow schemas whenever possible.

Reason:

- reliability
- monitoring
- fallback support

Principle 4 — Human Ownership Required

Humans approve:

- outreach
- CRM updates
- relationship decisions

AI Capability Inventory

Capability	AI Usage
Signal Interpretation	AI
Contextual Ranking Assistance	Hybrid
Recommendation Explanations	AI
Preparation Assistance	AI
Workflow Summaries	AI
Signal Collection	No AI
Workspace Rendering	No AI
Monitoring	No AI
Feedback Collection	No AI

AI System Specifications

AI System 1 — Signal Interpretation System

Purpose: Transform fragmented information into structured signals.

Inputs

- news
- job posts
- company websites
- announcements

Outputs: Structured Signal Objects.

Example:

Signal Type: Expansion

Confidence: Medium

Summary: Manufacturing expansion

Evidence: [...]

Why AI Here?

Interpretation across noisy, fragmented information becomes difficult using deterministic systems.

Non-AI Alternative

- rules
- keyword systems

- manual tagging

What Still Works Without AI?

- signal collection
- basic detection

What Is Lost?

- contextual interpretation
- adaptability
- higher recall

Autonomy Level

1. Interpret
2. Not Decide

AI System 2 — Contextual Ranking Assistance

Purpose: Provide ranking support using context signals.

Inputs

- signals
- relationship history
- previous activity
- business rules

Outputs

- ranking signals
- priority suggestions

Constraint

- Final ranking remains hybrid.
- AI contributes context.
- Rules enforce constraints.

Why AI Here?

Relationship context becomes difficult to encode entirely using rules.

Non-AI Alternative

Weighted scoring models.

What Still Works?

- ranking
- prioritization

What Is Lost?

- contextual understanding
- adaptability

Autonomy Level

1. Recommend
2. Not Decide

AI System 3 — Recommendation Explanation System

Purpose: Generate understandable explanations.

Inputs

- opportunities
- supporting signals

Outputs

- summary
- supporting signals
- confidence
- recommended action

Constraints

- summary ≤300 characters
- minimum 2 supporting signals

Why AI Here?

Summarization and explanation require synthesis.

Non-AI Alternative

Static templates.

What Is Lost?

- readability
- dynamic explanations

Autonomy Level

1. Explain
2. Not Decide

AI System 4 — Preparation Assistance System

Purpose: Reduce preparation effort.

Inputs

- signals
- CRM history
- relationship context

Outputs

- summaries
- conversation angles
- optional drafts

Constraints

- editable only
- no auto-send
- no auto-save

Why AI Here?

Requires synthesis across fragmented information.

Non-AI Alternative

- templates
- manual research

What Is Lost?

- preparation speed
- personalization

Autonomy Level

1. Prepare
2. Not Execute

AI System 5 — Workflow Assistance System

Purpose: Reduce administrative burden.

Inputs

- calls
- emails
- notes

Outputs

- summaries
- follow-up suggestions
- CRM suggestions

Constraints

- approval required
- automatic modification prohibited

Why AI Here?

Conversations are largely unstructured.

Non-AI Alternative

- templates
- manual notes

What Is Lost?

- productivity gains
- summarization benefits

Autonomy Level

1. Assist
2. Not Execute

Shared AI Constraints

All AI systems require:

- schema validation
- confidence fields
- traceability
- output storage

Confidence Handling

- High: Normal behavior
- Medium: Require additional review
- Low: Lower ranking weight
- Never: Trigger actions automatically

Failure & Fallback Systems

- Failure: Model unavailable or Inference Failure
  - Response: Deterministic fallback
- Failure: Low confidence
  - Response: Reduce recommendation ranking weight
- Failure: Invalid output schema
  - Response: Reject output
  - Trigger fallback
- Failure: High error rate detected
  - Response: Escalate
  - Rollback

Model Governance Requirements

System supports:

- model replacement
- prompt replacement
- versioning
- A/B testing

Reason: Models change. Evaluation standards remain.

AI Launch Requirements

AI capabilities cannot launch without:

- monitoring
- fallback systems
- rollback capability
- evaluation
- schema validation
- traceability

Evaluation Specification

This section defines how Linnk Atlas evaluates:

- recommendation quality
- trust
- degradation
- business impact

before launch and after launch.

Evaluation exists to answer: **Is Atlas improving decisions?** rather than: **Is the model still running?**

Evaluation is treated as product infrastructure rather than operational reporting.

Evaluation Principles

Principle 1 — Evaluation Before Models

- Tasks define evaluation.
- Evaluation defines acceptable behavior.
- Models are selected later.

Principle 2 — Production Evaluation Is Continuous

- Launch does not complete evaluation.

- Evaluation continues throughout product lifecycle.

Principle 3 — Business Outcomes Matter More Than Model Metrics

- Strong model metrics without behavioral change are insufficient.

Principle 4 — Trust Is Explicitly Measured

- Usage metrics and trust metrics are evaluated independently.

Evaluation Ownership

Product Team Owns

- evaluation definitions
- launch decisions
- thresholds

AI Engineering Owns

- evaluation execution
- monitoring infrastructure

Sales Own

- usefulness reviews
- human evaluation

Evaluation Framework

Linnk Atlas evaluation consists of:

- offline evaluation
- production evaluation
- human review
- drift monitoring
- launch / hold / rollback decisions

Evaluation Dataset Requirements

Required For:

- signal interpretation
- recommendation quality
- preparation assistance
- workflow summaries

Requirements:

- representative data
- recent data
- human reviewed samples
- versioned datasets

Datasets must support:

- multiple customers
- multiple industries
- multiple signal types

Requirement: Datasets remain reproducible.

Offline Evaluation

Purpose: Validate capabilities before launch.

Workflow:

1. Dataset Created
2. Run Evaluation
3. Human Review
4. Threshold Validation
5. Launch Recommendation

Requirements:

- versioned datasets
- human review
- approval required

Individual Task Evaluation Requirements

Signal Interpretation Evaluation

Questions: Did system identify signal correctly?

Metrics:

- Precision
- Recall
- Evidence Accuracy
- Human Rating

Dataset: Human labeled signals

- Launch Threshold:
    - Precision: 85%
- Rollback Trigger:
    - Precision: <70%
- Cadence: Weekly

Recommendation Explanation Evaluation

Questions:

- Is explanation understandable?
- Does explanation cite evidence?

Metrics:

- Evidence Presence
- Human Rating
- Acceptance Rate

Dataset: Sampled Recommendations

Launch Threshold:

- Evidence Presence: 100%
- Human Rating: ≥4/5
- Cadence: Weekly

Preparation Assistance Evaluation

Questions:

- Useful?

- Relevant?
- Correct?

Metrics:

- Human Rating
- Regeneration Rate
- Suggestion Usage

Dataset: Preparation Samples

Launch Threshold:

- Human Rating: ≥4/5
- Rollback Trigger:
  - Regeneration: 40%
- Cadence: Weekly

Production Evaluation

Purpose: Measure real-world recommendation quality.

Metrics:

Recommendation Quality Metrics

- recommendation acceptance rate
- dismissal rate
- meeting creation rate
- opportunity conversion rate

Trust Metrics

- override rate
- abandonment rate
- repeat usage
- feedback sentiment

Business Metrics

- research time reduction
- selling time increase
- meetings created
- pipeline growth
- revenue contribution

Traceability Requirements: Every recommendation stores:

- recommendation ID
- supporting signals
- confidence
- model version
- timestamps

Reason: Failures must be reconstructable.

Human Review Process

Frequency: Weekly

Review Inputs:

- random recommendations
- failure cases
- low confidence outputs

Outputs:

- approve
- investigate
- escalate
- rollback recommendation

Drift Monitoring

Monitor:

- signal drift
- behavior drift
- recommendation drift
- trust drift

Trigger: Sustained deterioration.

Action: Investigation required.

Rollback Framework

Example Triggers:

- dismissal rate >60%
- human ratings <3/5
- trust deterioration
- severe recommendation failures

Rollback Actions:

- reduce exposure
- deterministic fallback
- investigation
- recovery validation

Launch Gates

Launch blocked unless:

- offline evaluation passes
- monitoring operational
- feedback collection operational
- rollback mechanisms operational
- traceability operational
- human review operational

Trust & Failure Design

This section defines mechanisms ensuring Linnk Atlas remains:

- understandable
- trustworthy
- recoverable

when recommendations are incorrect, uncertain, incomplete, or unavailable.

Atlas is intentionally designed assuming failures will occur. Trust is therefore treated as product infrastructure rather than product output.

Trust Principles

- Principle 1 Evidence Before Recommendation
  - Users should understand why recommendation exists before acting on recommendation.
- Principle 2 Users Must Retain Control (decision ownership)
  - Atlas influences decisions. Atlas does not own decisions.
- Principle 3 Uncertainty Must Be Visible
  - Low confidence should not appear identical to high confidence.
- Principle 4 Failures Must Remain Recoverable
  - Users should recover quickly from recommendation failures.
- Principle 5 Trust Is Continuously Measured
  - Trust degradation must become visible.

Trust Mechanisms

Every displayed recommendation requires:

- explanation
- supporting evidence
- signal recency
- confidence indicator
- suggested action

Recommendation Display Rules

Recommendations:

- without explanations → not displayed
- without evidence → not displayed
- low confidence → excluded from Top 5

Every recommendation exposes:

- rationale
- evidence source
- timestamps

Confidence & Recommendation Behavior

High Confidence

Behavior:

- displayed normally
- all actions available

Medium Confidence

Behavior:

- review indicators displayed
- lower ranking weight

Low Confidence

Behavior:

- displayed as Potential Opportunity
- reduced visibility
- evidence emphasized

Highest Risk Feature Analysis

Highest Risk Feature: Opportunity Recommendations

Why Highest Risk?

This capability directly influences:

- who Sarah contacts
- where time is spent
- perceived recommendation quality

Example Failure

- Atlas recommends: Company A
- Actual opportunity exists at: Company B

Failure Impact

- wasted selling time
- trust degradation
- recommendation abandonment

Required Mitigations

- evidence visibility
- confidence indicators
- dismiss mechanisms
- feedback collection
- rollback capability

Failure Handling Framework

Signal Failures

Examples:

- missing signals
- delayed signals
- incomplete ingestion

User Impact: Weak recommendations

System Response:

- display freshness indicators
- fallback to cached signals
- lower recommendation confidence

Recommendation Failures

Examples:

- irrelevant opportunities
- poor ranking
- weak explanations

System Response:

- allow dismissal
- capture feedback
- reduce confidence weighting

Summarization Failures

Examples:

- inaccurate summaries

- missing context
- misleading preparation

System Response:

- edit
- regenerate
- reject

Reason: Users remain owners.

Integration Failures

Examples:

- CRM unavailable
- Signal provider unavailable
- ATS unavailable

System Response: Continue operating with partial functionality

Display: availability indicators

System Failures

Examples:

- Recommendation service unavailable
- Inference failures
- Latency failures

System Response: Fallback workflows.

Display: reduced capability states.

Failure Escalation Workflow

1. Failure Detected
2. Classified
3. Severity Assigned
4. Fallback Triggered
5. Investigation
6. Fix / Rollback / Recovery

Product Guardrails

Guardrails define explicit operational boundaries.

Atlas should:

- suggest opportunities
- suggest follow-ups
- generate preparation support
- suggest CRM updates

Atlas should Not:

- send communication automatically
- modify records automatically
- remove human approval
- own relationship decisions

Reason: Relationship ownership remains human.

Trust Monitoring

Monitor:

- recommendation acceptance
- dismissal rates
- override rates
- abandonment
- feedback sentiment

Trust considered healthy when:

- recommendations remain understandable
- users recover from failures
- recommendation usage remains stable
- trust does not deteriorate over time

Rollback Framework

Rollback Conditions:

- Trust decline
- Recommendation degradation
- Major failure incidents
- User abandonment increase

Rollback Actions:

1. Reduce traffic exposure
2. Disable capability
3. Fallback workflow
4. Recovery validation

Metrics & Analytics

This section defines metrics used to determine whether Linnk Atlas is:

- creating value
- improving decisions
- maintaining quality
- preserving trust

This section answers: **Is Atlas improving outbound sales decisions?** rather than: **Are users simply interacting with the product?**

Metrics Principles

- Principle 1 — Business Outcomes Matter More Than Engagement
  - Atlas exists to improve sales productivity rather than generate activity.
- Principle 2 — Usage Metrics And Quality Metrics Are Different
  - High usage does not imply high quality.
- Principle 3 — Trust Must Be Measured Explicitly
  - Recommendation quality alone is insufficient.
- Principle 4 — Quiet Abandonment Must Become Visible
  - Users silently abandoning recommendations should become measurable.

North Star Metric

Recommendation Influence Rate

- Definition: Percentage of outbound activities influenced by Atlas recommendations.
- Formula: Outbound activities initiated from Atlas / Total outbound activities

- Reason: Atlas exists to influence decisions. Not merely generate recommendations.
- Target: Increase continuously after launch.
- Declining trend for 4 consecutive weeks, investigation required

Business Metrics

These metrics measure whether Atlas creates commercial impact.

1. Research Time Reduction
   - Definition: Reduction in time spent researching accounts.
   - Formula: Baseline Research Time - Current Research Time
   - Reason: Research reduction is primary value proposition.
   - Alert: No measurable reduction after rollout.
2. Meetings Created
   - Definition: Meetings generated from Atlas influenced activities.
   - Purpose: Measures downstream business impact.
3. Selling Time Increase
   - Definition: Additional time spent on selling activities.
   - Purpose: Measure productivity improvement.
4. Pipeline Contribution
   - Definition: Pipeline value influenced by Atlas.
   - Purpose: Longer-term business measurement.

Behavior Metrics

Purpose: Measure whether users change behavior.

1. Workspace Open Rate
   - Definition: Percentage of active users opening Atlas.
   - Reason: Atlas should become morning workflow.
2. Time To First Action
   - Definition: Time between workspace load and first action.
   - Purpose: Measures decision friction.
3. Recommendation Interaction Rate
   - Definition: Recommendations interacted with.
   - Purpose: Detect ignored recommendations.
4. Preparation Usage Rate
   - Definition: Usage frequency of preparation workflows.

Quality Metrics

Purpose: Measure recommendation quality.

1. Recommendation Acceptance Rate
   - Definition: Recommendations resulting in user action.
   - Target: ≥30%
   - Alert: Sustained decline.
2. Recommendation Dismissal Rate
   - Definition: Recommendations explicitly dismissed.
   - Purpose: Detect poor quality.
   - Alert Threshold: 60%
   - Action: Investigation required.
3. Regeneration Rate
   - Definition: Frequency generated outputs require regeneration.
   - Purpose: Measure output usefulness.
4. Summary Edit Rate
   - Definition: Extent generated summaries require modification.

Trust Metrics

Purpose: Measure trust directly.

1. Override Rate
    - Definition: Recommendations ignored or replaced.
    - Reason: High overrides indicate trust problems.
2. Recommendation Confidence Calibration
    - Question: Do high confidence recommendations outperform lower confidence recommendations?
    - Reason: Confidence should mean something.
3. Workspace Abandonment
    - Definition: Users opening workspace and leaving without action.
    - Reason: Detect silent failure.
4. Feedback Sentiment
    - Definition: Qualitative user feedback trends.

Operational Metrics

Purpose: Measure reliability.

1. Workspace Load Time
    - Target: <5 sec P95
2. Recommendation Generation Latency
    - Target: <5 sec P95
3. Availability
    - Target: 99.5%
4. Signal Freshness
    - Target:
        - CRM: ≤15 mins
        - External: ≤24 hrs

Review Cadence

1. Operational Metrics: Continuous
2. Quality Metrics: Weekly
3. Trust Metrics: Weekly
4. Business Metrics: Monthly

Rollout Strategy

This section defines how Linnk Atlas AI transitions from validation to scaled deployment while minimizing:

- trust failures
- workflow disruption
- operational risk

Atlas changes sales behavior. Behavior-changing systems require gradual rollout rather than large launches.

Rollout Principles

1. Integrate Before Replace
    a. Atlas launches alongside CRM, ATS, and existing workflows
    b. Reason: Replacement creates adoption risk.
2. Capabilities Roll Out Before Automation
    a. Trust should be earned before expanding autonomy.
3. Expansion Requires Evidence
    a. Progression between phases requires measurable outcomes rather than timelines..
4. Rollout Must Remain Reversible: Every rollout stage supports:
    a. feature disablement
    b. traffic reduction
    c. fallback workflows

Rollout Model

Atlas rollout occurs across four phases.

Phase 0 — Internal Validation

Objective: Validate product functionality.

Participants: Internal Teams

Activities:

- integration testing
- recommendation testing
- workflow validation
- evaluation validation

Capabilities Enabled:

- Signal Processing
- Opportunity Generation
- Workspace
- Preparation Assistance

Success Requirements:

- Core workflows functional
- Monitoring operational
- Evaluation operational
- Rollback tested

Exit Criteria:

- Evaluation thresholds passed
- No major workflow blockers

Phase 1 — Pilot Agencies

Objective: Validate usefulness with real users.

Target Customers: Small number of agencies.

Recommended: 5–10 agencies.

Recommended Users: High activity users.

Capabilities Enabled:

- Opportunity Intelligence
- Workspace
- Preparation Assistance
- Limited Workflow Assistance

Capabilities Disabled:

- Broad automation
- Expanded integrations
- Advanced workflows

Pilot Duration: 4–8 weeks

Primary Questions:

1. Do recommendations influence behavior?
2. Do users trust outputs?
3. Do users open Atlas daily?

Exit Criteria:

- Recommendation acceptance healthy
- Trust metrics healthy
- Research reduction measurable
- No severe trust incidents

Phase 2 — Controlled Rollout

Objective: Expand usage while controlling risk.

Participants: Larger customer set.

Capabilities Added:

- Workflow Assistance
- Expanded Integrations
- Improved Personalization

Rollout Mechanism:

- Feature Flags
- Exposure Controls
- Capability Toggles
- Monitoring Requirements:
- Weekly reviews mandatory
- Drift monitoring active
- Rollback tested

Exit Criteria:

- Stable quality
- Stable trust
- Operational readiness

Phase 3 — Scaled Deployment

Objective: Broad availability.

Capabilities Enabled: Full MVP capability set.

Requirements:

- Monitoring active
- Evaluation active
- Rollback active
- Support processes operational
- Rollout remains reversible.

Capability Rollout Sequencing

1. Opportunity Intelligence, this is core value proposition.
2. Linnk Workspace, makes intelligence actionable.
3. Preparation Assistance, reduces research burden.
4. Workflow Assistance, Higher operational risk.

Principle: **Decision Support Before Operational Assistance**

Launch Gates & Rollback Strategy

Capabilities cannot progress unless:

- evaluation passes
- trust remains healthy
- monitoring operational
- rollback validated

Rollback Triggers:

- trust deterioration
- recommendation failures
- quality degradation

Rollback Actions:

- reduce exposure
- disable capabilities
- activate fallback workflows
- every rollout must support rollback

Risks & Mitigations

This section identifies major riskss that can prevent Linnl Atlas from:

- improving decisions
- earning trust
- scaling safely

Atlas assumes Failures will occur. This section therefore defines:

- what can fail
- how failure becomes visible
- how failure is contained

This section exists because AI systems fail differently from deterministic software. Risk management therefore becomes product work. Not post-launch work.

Risk Framework

- Risk severity determined by: **Likelihood × Impact**
- Risk Levels: Critical, High, Medium, Low
- Critical risks require:
  - Active monitoring
  - mitigation plans
  - rollback plans

Product Risks

Recommendation Quality Risk

1. Severity: Critical
2. Description: Recommendations consistently surface poor opportunities or recommendation quality gradually deteriorates.
3. Business Impact
   a. wasted selling effort
   b. reduced recommendation influence
   c. trust deterioration
4. Leading Indicators
   a. increasing dismissals
   b. increasing overrides
   c. declining acceptance
   d. declining influence rates
5. Mitigations
   a. Increase evidence visibility
   b. Reduce exposure
   c. Increase review requirements
   d. Fallback ranking

Quiet Trust Degradation

1. Severity: Critical
2. Description: Users continue opening Atlas but stop relying on recommendations.
3. Business Impact: Product appears healthy while value collapses.
4. Leading Indicators
   a. interactions decline
   b. abandonment increases
   c. influence decreases
5. Mitigations
   a. Increase feedback collection
   b. Review explanations

    c. Reduce weak recommendations
    d. Customer interviews

Workflow Resistance

1. Severity: Critical
2. Description: Users continue CRM-first workflows.
3. Business Impact:
    a. Low adoption.
    b. Low influence.
4. Leading Indicators
    a. workspace usage weak
    b. preparation usage weak
    c. influence rates low
5. Mitigations
    a. Gradual rollout
    b. Workflow integration
    c. Customer success support
    d. Behavioral onboarding

AI Risks

Incorrect Summaries / Hallucinations

1. Severity: High
2. Description: Generated outputs contain incorrect information.
3. Business Impact
    a. poor preparation
    b. reduced trust
4. Leading Indicators
    a. edit rates increase
    b. regeneration increases
5. Mitigations
    a. schema validation
    b. evidence constraints
    c. approval workflows

Over-Autonomy Risk

1. Severity: Critical
2. Description: Automation expands faster than trust.
3. Business Impact:
    a. Reputation damage.
    b. Operational failures.
4. Mitigations
    a. approval workflows
    b. feature flags
    c. autonomy reviews
    d. capability gating

Technical Risks

Weak Signal Quality

1. Severity: High
2. Description: Signals become stale, noisy, or incomplete.
3. Business Impact:
    a. Poor recommendations.
    b. Lower trust.
4. Leading Indicators:
    a. Signal freshness deterioration
    b. Coverage reduction
    c. Recommendation quality decline
5. Mitigations

    a. Multiple signal providers
    b. Caching
    c. Confidence reduction
    d. Fallback workflows

Integration Complexity

1. Severity: High
2. Description: CRM / ATS integrations become difficult.
3. Business Impact:
    a. Delays.
    b. Support burden.
    c. Data issues.
4. Mitigations
    a. modular integrations
    b. limited initial integrations
    c. feature flags

Commercial Risks

Unit Economics Risk

1. Severity: Medium
2. Description: AI costs scale faster than value.
3. Business Impact: Weak commercial viability.
4. Mitigations
    a. Caching
    b. Hybrid systems
    c. Asynchronous processing
    d. Reduce expensive inference paths

Risk Management Approach

Linnk Atlas manages risk through:

- monitoring
- evaluation
- rollback mechanisms
- feature flags
- controlled rollout

Risk management succeeds when:

- failures become visible early
- trust remains stable
- risks remain recoverable

Dependencies & Assumptions

This section defines dependencies and assumptions required for:

- product success
- rollout feasibility
- implementation realism

Atlas operates as an intelligence layer above existing workflows. Its success therefore depends on systems, data, and behaviors outside Atlas itself.

Atlas assumes:

- dependencies may fail
- customer environments remain imperfect
- fallback mechanisms remain necessary

Dependencies should therefore be observable, manageable and replaceable where possible.

External Dependencies

CRM Availability

1. Description: Atlas requires access to existing CRM systems.
2. Dependency Type: Critical
3. Used For:
    a. Relationship Context
    b. Previous Activity
    c. Account Ownership
4. Constraint: CRM remains system of record.
5. Failure Impact: Reduced recommendation quality.
6. Assumption: Customers allow read integrations.

ATS Availability

1. Description: Atlas requires ATS integration.
2. Used For:
    a. Hiring Context
    b. Previous Placements
    c. Customer Activity
3. Failure Impact: Reduced signal quality.
4. Assumption: Customers maintain usable ATS data.

External Signal Providers

1. Description: Atlas depends on external signal sources.
2. Examples:
    a. Job Boards
    b. Company Websites
    c. Hiring Pages
    d. News Sources
    e. Tenders
3. Failure Impact: Signal quality degradation.
4. Assumption: Multiple providers available.

Data Dependencies

Data Quality

1. Description: Recommendations depend on:
    a. accurate records
    b. usable activity history
    c. account ownership accuracy
2. Failure Impact: Weak personalization.
3. Assumption: Customer data quality is imperfect.
4. Product Response: Operate under incomplete information.

Historical Context Availability

1. Description: Certain capabilities require historical context.
2. Examples:
    a. Preparation Assistance
    b. Relationship Context
    c. Workflow Assistance
3. Assumption: Historical data exists for at least some accounts.
4. Constraint: Product must operate with limited history.

Operational Dependencies

Monitoring Infrastructure

1. Required For:
    a. quality monitoring
    b. trust monitoring
    c. rollback support

2. Constraint: Capabilities cannot launch without monitoring.

Feedback Infrastructure

1. Required For:
    a. evaluation
    b. recommendation improvement
    c. trust measurement
2. Constraint: Feedback infrastructure treated as mandatory.

Human Review Availability

1. Required For:
    a. evaluation reviews
    b. failure investigation
    c. recommendation validation
2. Assumption: Sales SMEs remain available.

Product Assumptions

1. Users Prefer Assistance Over Automation
    - Reason: Staffing relationships remain human-driven.
    - Impact: Autonomy intentionally limited.
2. Recommendation Quality Matters More Than Recommendation Quantity
    - Reason: Decision reduction is primary value.
    - Impact: Maximum opportunities intentionally limited.
3. Existing Workflows Will Persist
    - Reason: Customers already use CRM, ATS, Spreadsheets, Email workflows etc.
    - Impact: Atlas designed for coexistence.
4. Trust Is Harder To Earn Than Functionality
    - Reason: Recommendation products fail when trust fails.
    - Impact: Evidence and confidence required.
5. Users Will Tolerate Imperfect Recommendations If Recovery Is Easy
    - Reason: Probabilistic systems produce mistakes.
    - Impact: Recovery mechanisms prioritized.

Explicit Non-Assumptions

Atlas does not assume:

- perfect CRM data
- perfect ATS data
- complete signals
- immediate adoption
- standardized workflows

Reason: Atlas designed for imperfect environments.

Open Questions & Deliberate Non-Goals

Objective

Capture unresolved questions, explicit boundaries, and deliberate decisions that shape Atlas beyond MVP.

This section exists because:

- certain decisions require real-world validation
- scope expansion requires discipline
- product boundaries should remain explicit

Open Questions Requiring Validation

These questions remain intentionally unresolved and require validation through rollout.

Q1. Which Signal Sources Create Highest Recommendation Quality?

1. Current Position: Atlas initially supports
    a. CRM signals
    b. ATS signals
    c. job boards
    d. company websites
    e. news
    f. hiring pages
2. Unknown: Relative contribution of each source.
3. Why This Matters: Signal acquisition increases
    a. complexity
    b. cost
    c. maintenance burden
4. Validation Plan: Track
    a. Recommendation Quality
    b. Signal Usage
    c. Recommendation Outcomes
5. Decision Trigger: Low value signals may be removed.

Q2. What Recommendation Volume Maximizes Decision Quality?

1. Current Position: Maximum: 20 opportunities
2. Unknown: How many opportunities creates best outcomes.
3. Validation Plan: Measure
    a. Acceptance Rate
    b. Decision Speed
    c. Influence Rate
4. Reason: More recommendations may reduce usefulness.

Q3. What Is The Long-Term Autonomy Ceiling?

1. Current Position: Assistive product.
2. Unknown: Whether users eventually want:
    a. More automation OR
    b. Stronger control.
3. Reason: Autonomy should be earned rather than assumed.

Deliberate Non-Goals

These decisions are intentional.

Fully Autonomous Outreach

1. Decision: Rejected.
2. Meaning: Atlas will NOT
    a. send emails automatically
    b. send messages automatically
    c. initiate communication automatically
3. Reason: Creates
    a. trust risk
    b. compliance risk
    c. relationship risk
4. Product Position: Atlas prepares. Humans communicate.

CRM Replacement

1. Decision: Rejected.
2. Meaning: Atlas will NOT become
    a. system of record
    b. full CRM platform
    c. operational workflow platform
3. Reason: Customers already have
    a. CRMs
    b. ATS systems
    c. Established workflows
4. Replacement increases:
    a. complexity

    b. adoption risk
    c. implementation burden
5. Product Position: Atlas is an intelligence layer not a replacement platform.

Exclusions From MVP

- candidate matching
- recruitment operations
- payroll / compensation workflows
- back-office automation

Reason: These capabilities do not directly improve: **Who should Sarah contact today?**

Future Development Decision Framework

Future capabilities require:

1. Clear problem definition
2. Deterministic alternatives evaluated
3. AI justification established
4. Trust impact acceptable
5. Business value demonstrated

Reason: Capability expansion should require evidence.

Linnk Atlas AI Product Prototype Screens

Please explore the interactive prototype for workflow demonstration and product visualization by clicking here: Linnk Atlas v.1.0 - Lead Generation & Staffing Sales

Note: This prototype is intended for illustrative purposes and does not represent final production implementation.

Screen 1: Linnk Workspace & Recommended Prospects Queue

- To help Sarah quickly decide who to call, when, and what to say when they pick up..
- Sarah Can:
    - View prioritized prospects with supporting signals and rationale
    - Understand why specific opportunities are recommended
    - Initiate outreach preparation directly from surfaced opportunities
    - Save, dismiss, or provide feedback on recommendations
    - Review follow-up reminders, overdue actions, and re-engagement opportunities

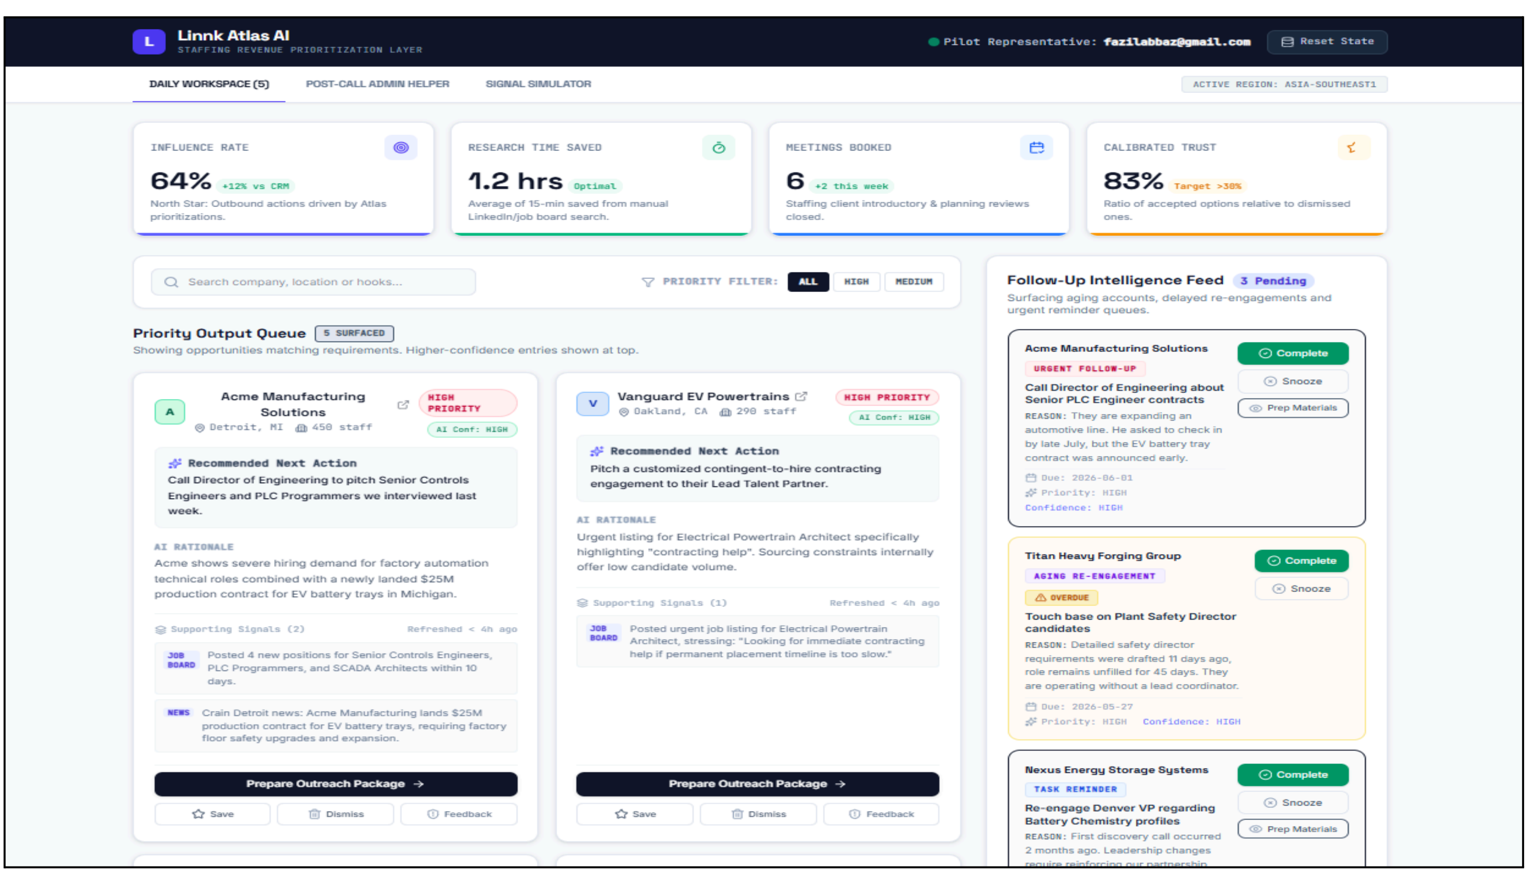

Screen 2: Outreach Preparation Intelligence

- To reduce Sarah's research effort before calls, emails, and outreach activities.
- Sarah Can:
    - View synthesized account context, hiring signals, and relationship history
    - Review recommended conversation hooks and outreach angles
    - Generate customizable outreach assets (email drafts, InMail drafts, messaging support, etc.)
    - Use AI-generated preparation packages while retaining full ownership of communication
    - Move from opportunity identification to outreach readiness within minutes

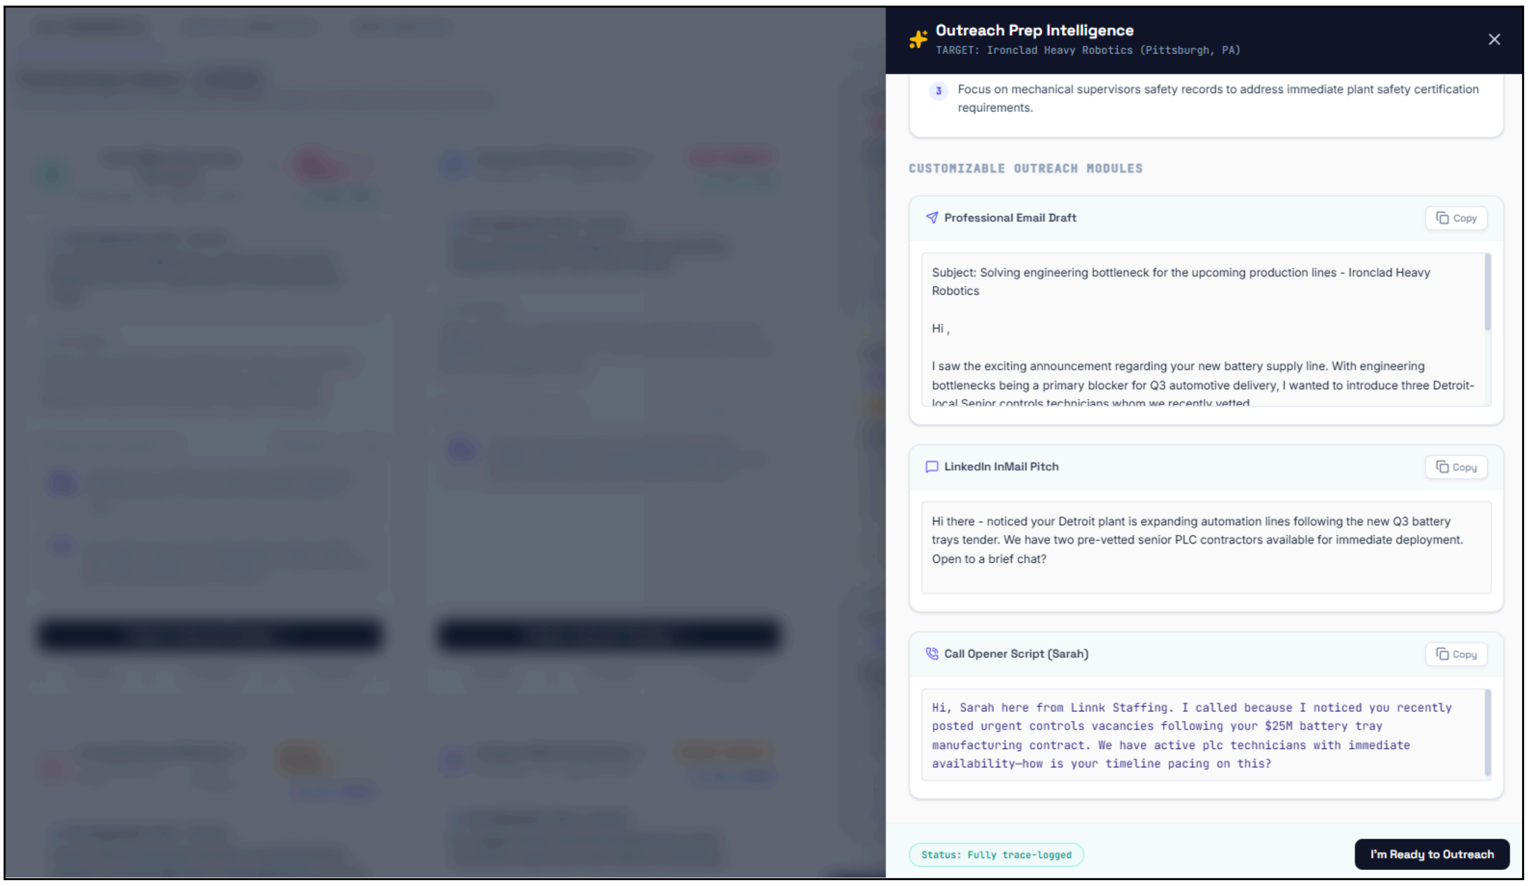

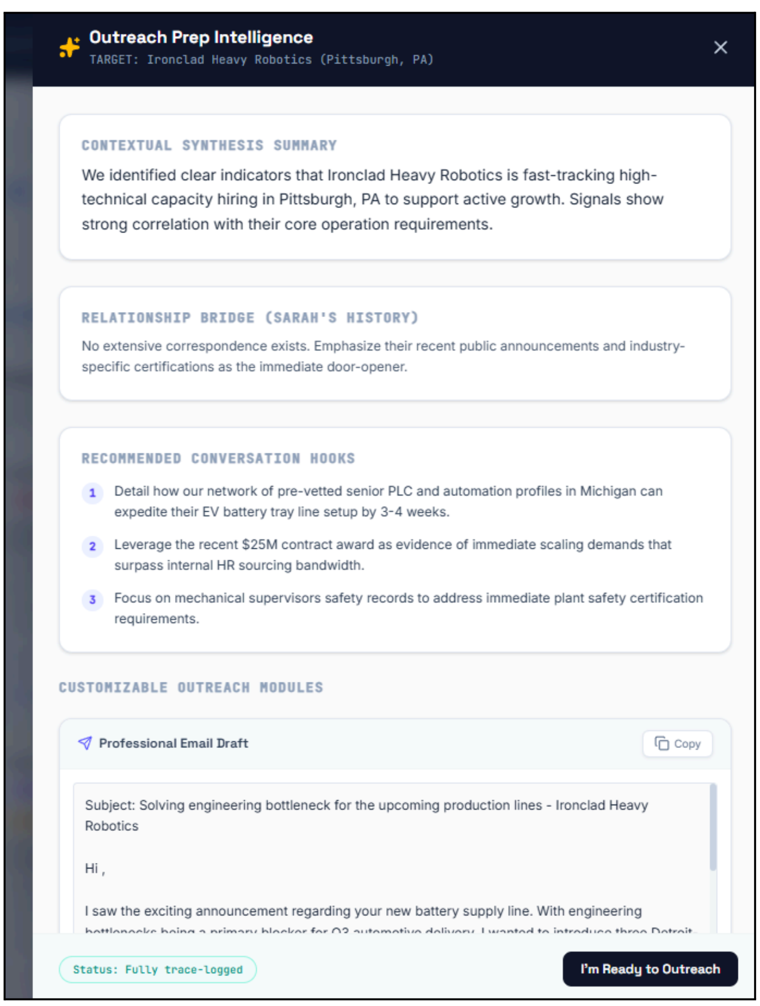

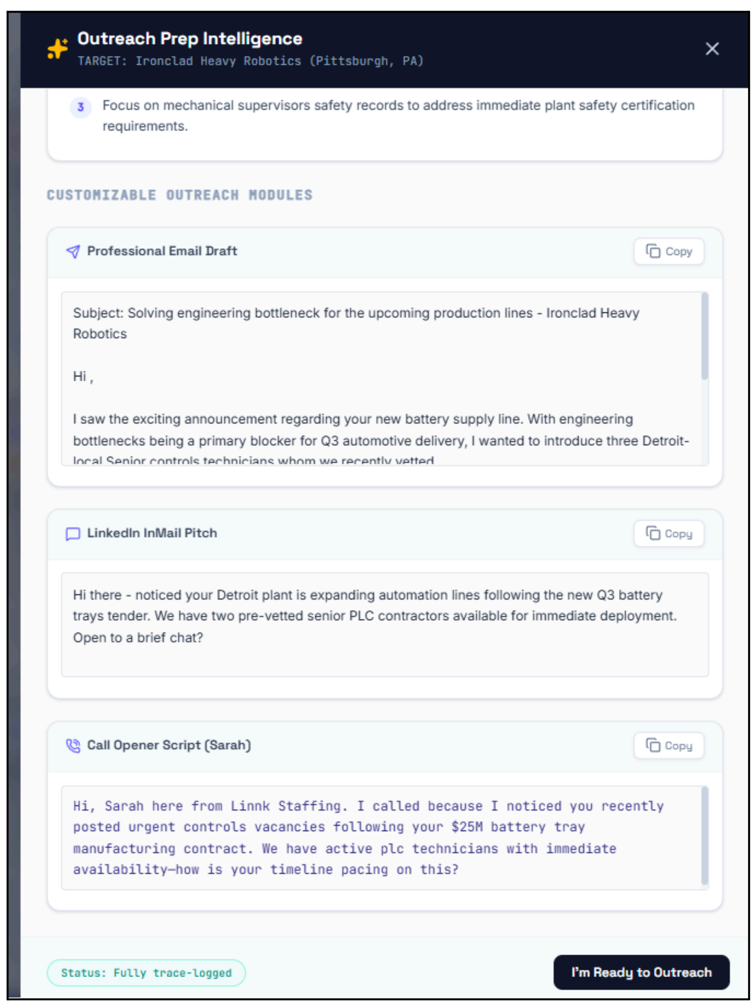

Screen 3: Follow-Up Intelligence Feed

- To help Sarah reduce missed opportunities by surfacing follow-ups, reminders, and re-engagement opportunities proactively.
- Sarah Can:

- View prioritized follow-up recommendations and overdue actions
- Understand why specific follow-ups are being surfaced
- Prepare follow-up materials and outreach support directly from recommendations
- Snooze, complete, or act on reminders based on current priorities
- Maintain timely engagement without manually tracking follow-ups across multiple systems

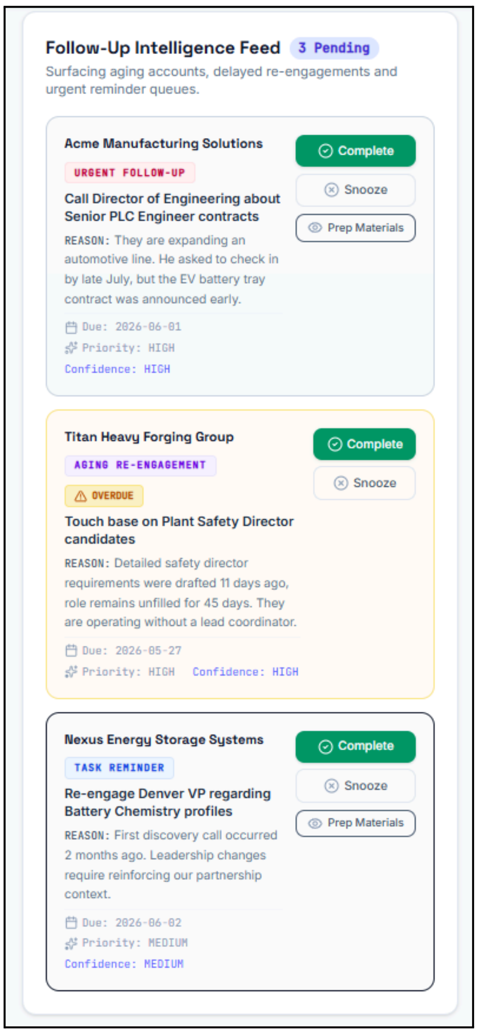

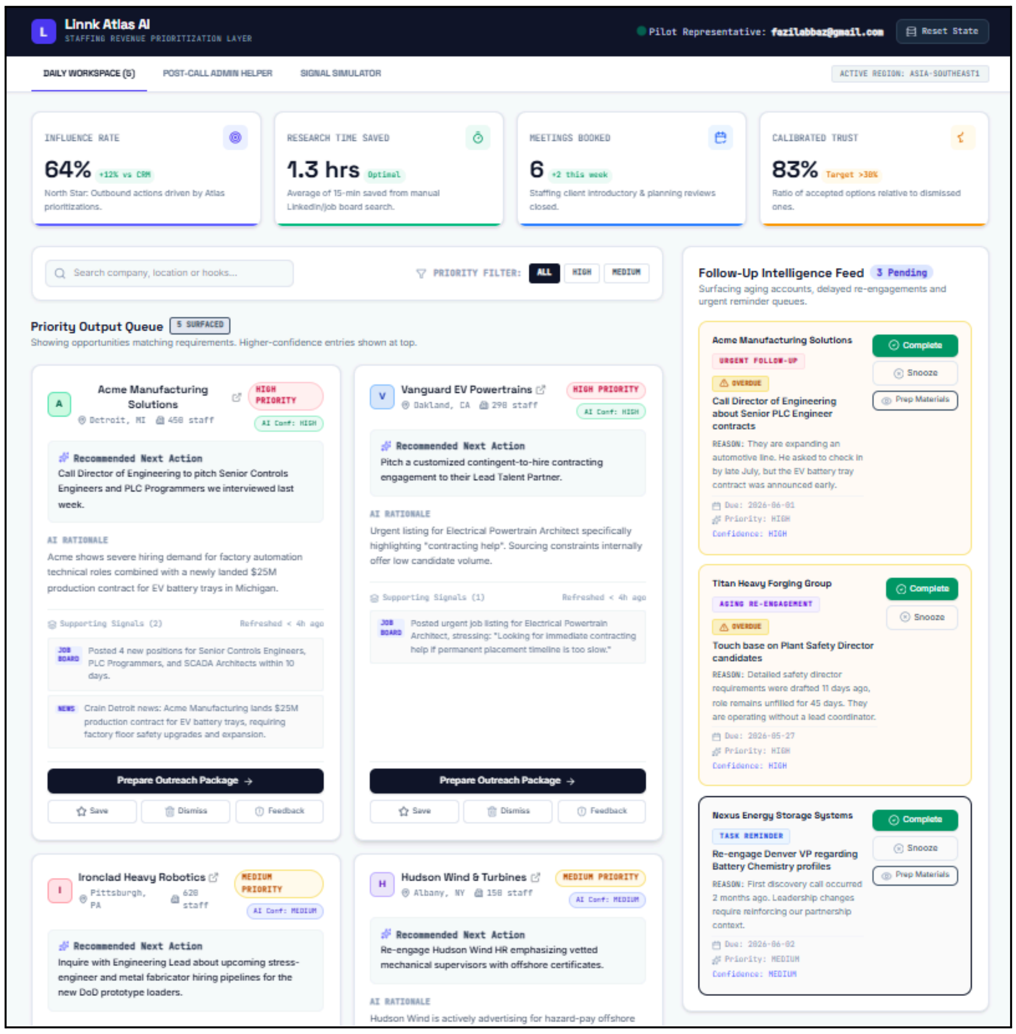

Screen 4: Post-Call Admin Helper

- To help Sarah reduce post-outreach administrative effort and ensure follow-ups are not missed.
- Sarah can:
  - Input rough interaction notes from calls, meetings, or emails
  - Generate structured summaries, CRM updates, and activity records from unstructured notes
  - Receive suggested follow-up actions, reminders, and next-step recommendations
  - Review, edit, approve, or reject generated outputs before CRM / ATS synchronization
  - Reduce manual documentation effort while retaining ownership of final decisions

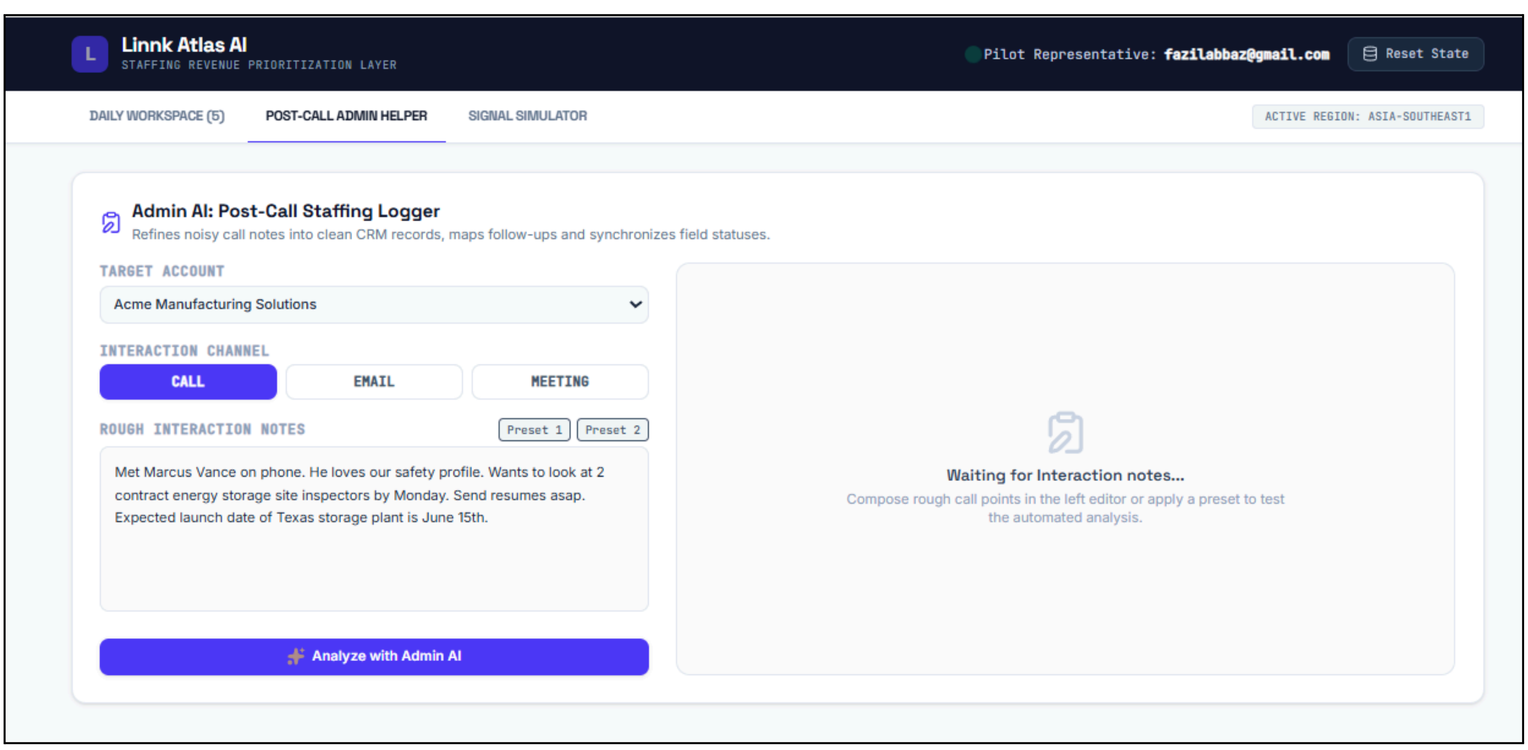

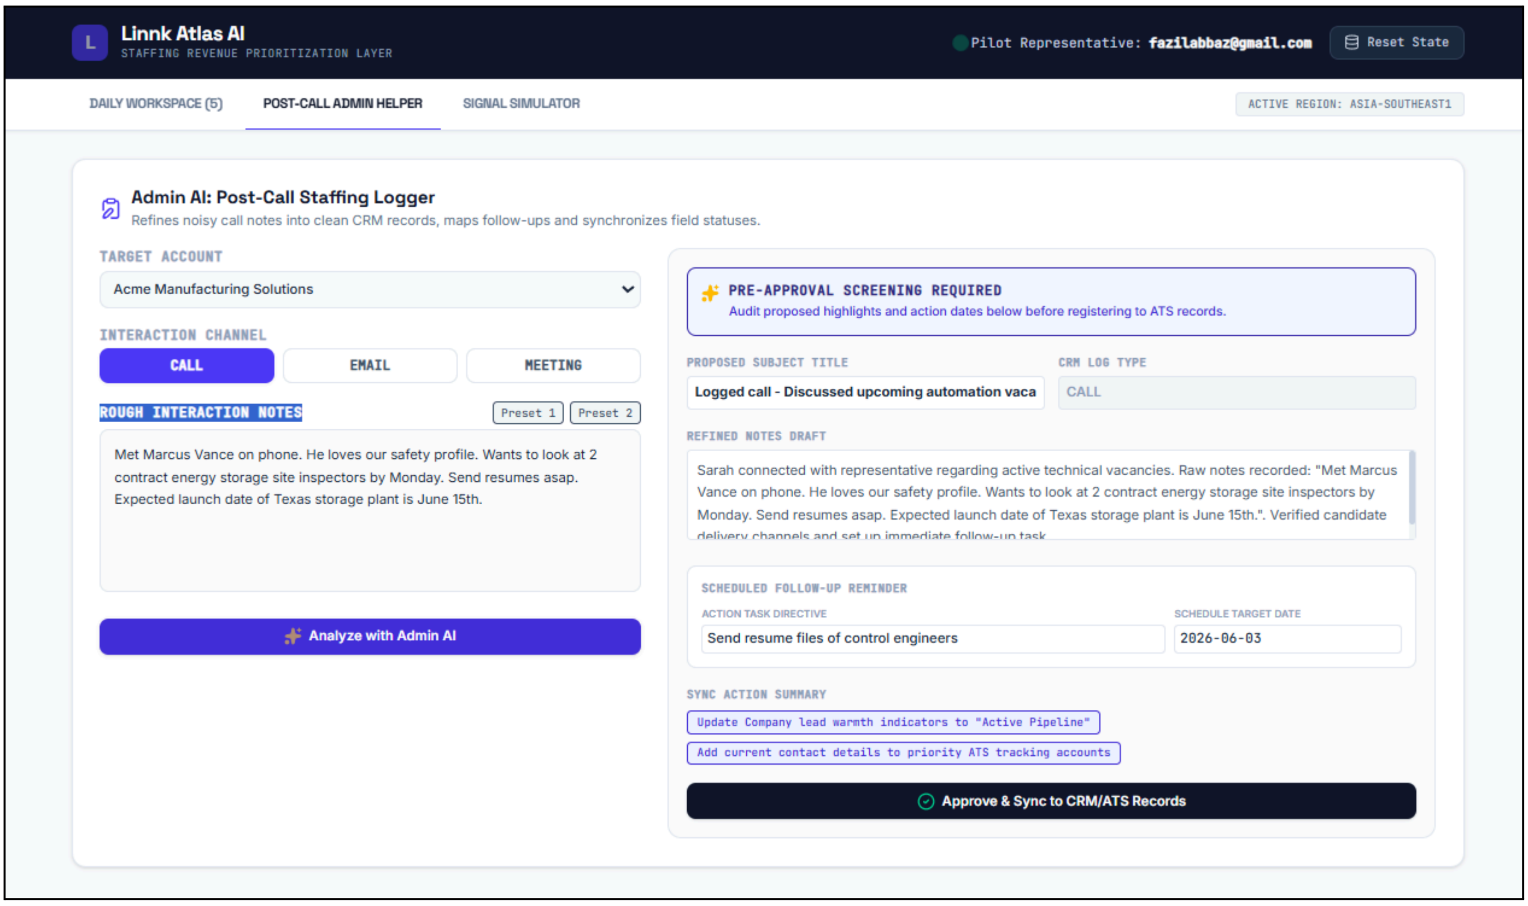

Screen 5: Market Signal Processing & Recommended Prospect Creation

- To Demonstrate how external market signals are transformed into actionable sales opportunities.
    - When external signals such as announcements, news updates, hiring activity, or market events are fed into the system.
    - Linnk Atlas AI interpret unstructured information and evaluate opportunity potential
    - Automatically populate qualified opportunities into the Daily Workspace when minimum signal requirements are satisfied
- Note: This workflow primarily operates in the background. This screen is included to demonstrate how recommended prospect generation occurs before recommendations appear within the daily workflow.

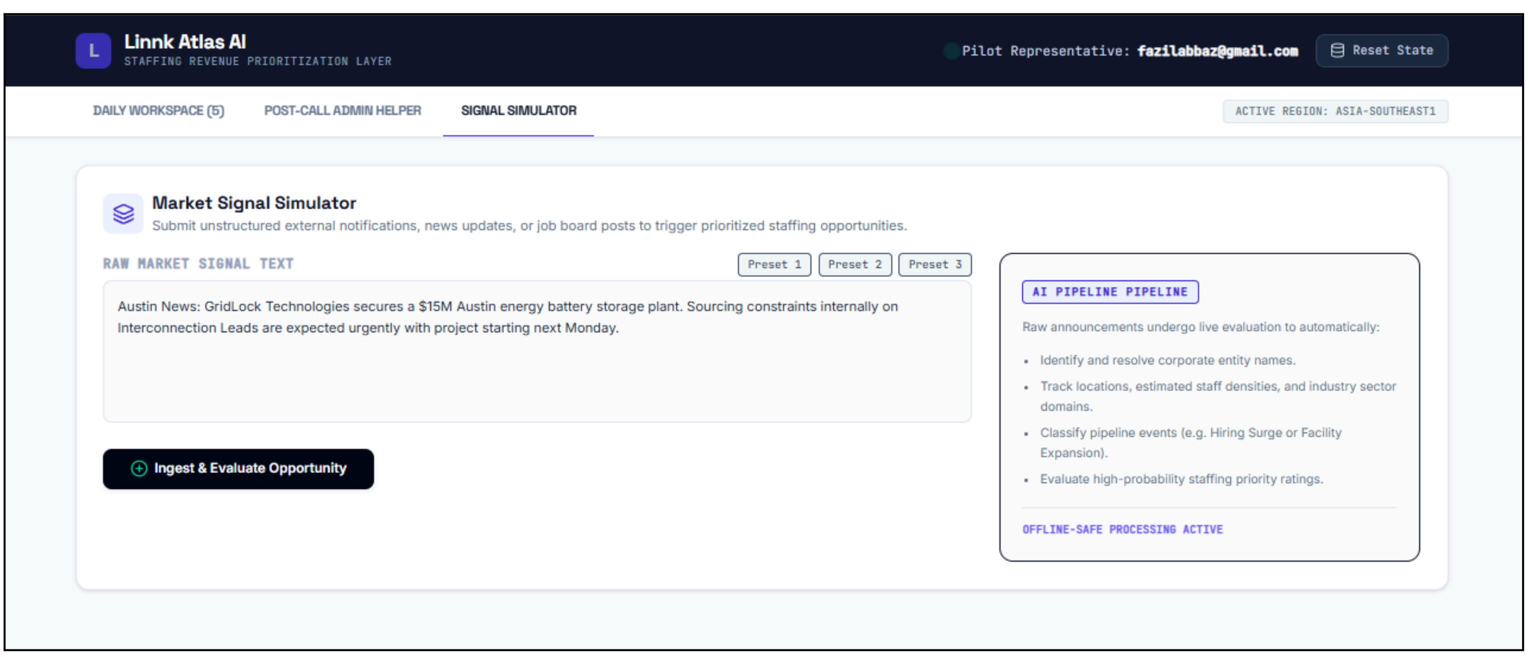

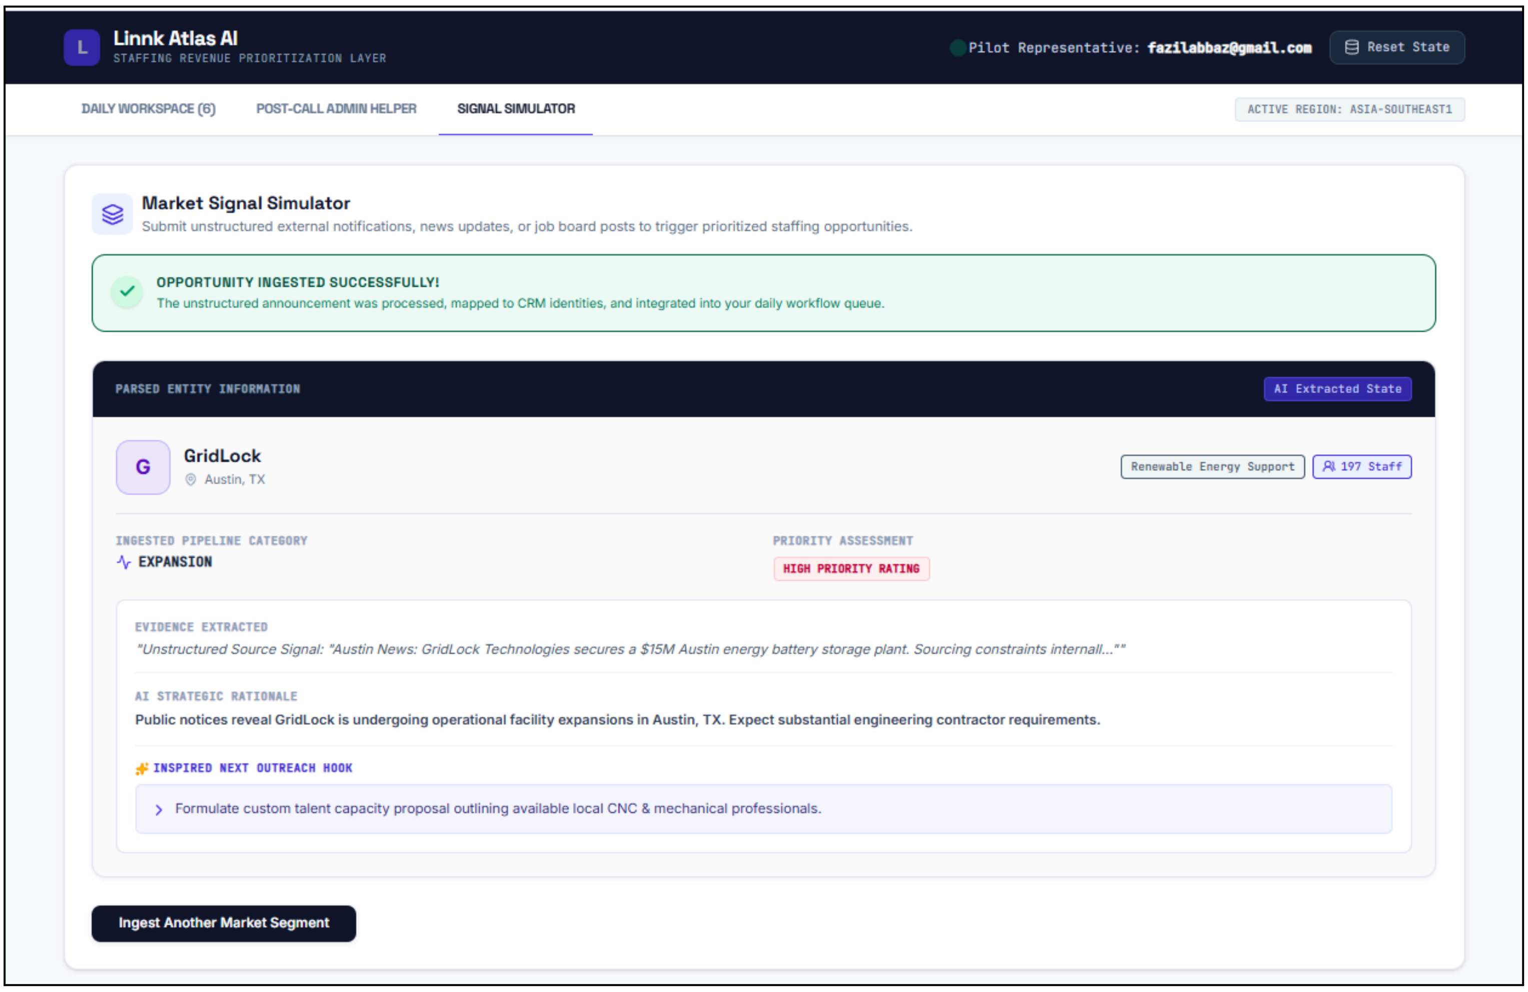

Closing Note

Linnk Atlas AI is intentionally designed as an intelligence layer that improves outbound decision quality for staffing sales teams. This PRD intentionally prioritizes product decisions, trade-offs, implementation clarity, and operational realism. Throughout this document, assumptions, constraints, and trade-offs have been explicitly documented, while scope has been intentionally constrained around solving a single primary problem: **helping Sarah know who deserves attention before outreach begins.**

Several product decisions intentionally reinforce this focus, including designing **Linnk Atlas AI** as an intelligence layer rather than a CRM replacement, limiting autonomy despite higher automation being technically possible, prioritizing decision quality before workflow optimization, and treating trust, evaluation, monitoring, and recovery mechanisms as product requirements.

Supporting prototype screens have been included in **Linnk Atlas AI Product Prototype Screens** section to visualize key workflows and interaction patterns. An **interactive prototype link has also been attached** to demonstrate how major product workflows may translate into user experiences and operational flows.

Traditional user story formats were intentionally not used. Instead, requirements and user flows are optimized primarily for implementation clarity, workflow context, and measurable requirements rather than strict adherence to conventional templates.

This PRD should be viewed as a planning artifact. Certain assumptions, thresholds, signal sources, and rollout decisions require validation through real-world usage. The objective of this document is to create sufficient clarity for Product, Design, Engineering, and Commercial teams to begin building, evaluating, and challenging decisions immediately.